

A Historical Dataset of U.S. Governors

Sydney White

Theo Serlin

Neel Gupta

Kenneth Scheve

2026-02-01

Data Essay

Introduction

This dataset contains information on the biographical backgrounds of United States governors (of states and territories) from 1775 to the present. You can explore and download the expanded data in the “Explore the Data” tab.

Brief Survey

If you use our materials in your class or another setting, we would love to [hear about it!](#)

Dataset

```
//|echo: false
// Using the Post45 Data Viewer
// https://observablehq.com/d/8bb63a6cde9addff
import {viewof dataSummaryView, Tabulator, viewof selectedColumns, viewof dataSet, tableCont
```

```
//|echo: false
//|output: false

// Using the Post45 Data Viewer
// https://observablehq.com/d/8bb63a6cde9addff
generateTabulatorTableFromCSV(
  "#table-container4",
  "https://raw.githubusercontent.com/melaniewalsh/responsible-datasets-in-context/refs/heads/
  {
```

```

displayedColumns: [
  // Core identity & office
  "state_territory",
  "governor",
  "gender",
  "race/ethnicity",

  "party",
  "first_year",
  "years_in_office",

  // Background & education
  "school",
  "college_attendance",
  "ivy_attendance",
  "lawyer",
  "military_service",

  // Birth & demographics
  "birth_state_territory",
  "birth_date",
  "age_at_start",

  "born_in_state_territory",
  "intl_born",
  "intl_born_details",

  // Long description
],

columnPopups: [
  "U.S. state or territory governed.",
  "Governor's name.",
  "Gender based on research.",
  "Race/ethnicity based on research.",
  "Political party affiliation.",
  "First year they served as governor.",
  "Spans/terms served (as listed in the bio).",

  "Primary college or university listed (if any).",
  "Did the governor attend college? (1/0).",
  "Did the governor attend an Ivy League school? (1/0).",

```

```

    "Is/was the governor a lawyer? (1/0).",
    "Did the governor serve in the military? (1/0).",

    "Birth state or territory.",
    "Birth date (YYYY-MM-DD where available).",
    "Age at the start of first term.",
    "Whether born in the state/territory they govern (1/0).",
    "Whether born outside the U.S. (1/0).",
    "Details about international birthplace if applicable.",

],

// columnWidths: {
//   governor: "220px",
//   party: "140px",
//   years_in_office: "220px",
//   school: "220px",
//   bio_text: "40%"
// },

// Treat these as numeric so you get proper sort/filter ranges
numericColumns: ["first_year", "age_at_start"],

// Good facets for filtering
categoryColumns: [
  "state_territory",
  "party",
  "school",
  "birth_state_territory",
  "gender",
  "race/ethnicity",

  // binary flags as categories for easy faceting
  "college_attendance",
  "ivy_attendance",
  "lawyer",
  "military_service",
  "born_in_state_territory",
  "intl_born"
],

// Default sort: by state, then governor

```

```

    sortColumns: ["state_territory", "governor"],
    sortOrders: ["asc", "asc"],

    // Buttons (kept the same IDs you used)
    buttonContainerId: "#button-container1",
    rawButtonId: "#download-raw1",
    urlCopyButtonId: "#copy-url1",

    // Optional: if your helper supports label maps for 1/0 columns, this is handy.
    // booleanLabelColumns: {
    //   college_attendance: { "1": "Yes", "0": "No" },
    //   ivy_attendance: { "1": "Yes", "0": "No" },
    //   lawyer: { "1": "Yes", "0": "No" },
    //   military_service: { "1": "Yes", "0": "No" },
    //   born_in_state_territory: { "1": "Yes", "0": "No" },
    //   intl_born: { "1": "Yes", "0": "No" }
    // }
  }
);

```

Download Full Data

Download Full Data

Copy Full Data URL

Download Table Data (including filtered options)

Download CSV

Download JSON

Download Excel

Creative Commons License

This work is licensed under CC BY 4.0

What's in the data?

Alongside the names and states of each governor, the dataset includes information on gender, educational background, as well as other characteristics. The dataset is based on the publicly-available biographies of governors from the [National Governors Association](#) (NGA). We capture the following variables from the website:

- `state_territory`: U.S. state or territory where the governor served.
- `governor`: Full name of the governor, as listed on the website. The listed name is not always the legal name of the person, e.g. John Ellis Bush Jr. is listed as Jeb Bush.
- `party`: Political party affiliation under which the governor served.
- `years_in_office`: Concatenated string denoting term(s) in office (e.g., “1985–1991; 1995–1999”). Parsed from web biographies. - `school`: Educational institution(s) attended by the governor as listed on the official profile. If multiple schools are listed, usually separated by semicolons (e.g., “Brigham Young University; Harvard University”). If NA, either no school is listed or the governor did not attend college.
- `birth_state_territory`: Parsed location (state) of birth based on text extraction.
- `spouse`: Name of spouse, if available. If NA, the governor is either not married or spousal information is unknown.
- `birth_date`: Date of birth, as listed on website.
- `bio_text`: Raw string of the biographical background of the governor. In 14 cases (less than 1 percent of cases), this biographical text is missing.

Using the above information extracted from the NGA, we construct the following columns:

- `college_attendance`: Equals 1 if the governor attended any college or university; 0 otherwise.
- `ivy_attendance`: Equals 1 if the governor attended an Ivy League institution; 0 otherwise. (Note that this technically applies for EITHER graduate or undergraduate).
- `lawyer`: Equals 1 if the biography indicates legal education or professional law practice; 0 otherwise.
- `military_service`: Equals 1 if the biography mentions military service (e.g., Army, Navy, National Guard, etc.); 0 otherwise.
- `age_at_start`: Age of the governor at the beginning of their first term, calculated from the reported date of birth and the start date of their first gubernatorial term.
- `gender`: We manually fill in gender based on the [Wikipedia list of female governors](#). All other governors are set to male.
- `born_in_state_territory`: 1 if the governor was born in the same U.S. state/territory they governed; 0 if born elsewhere (including outside the U.S.).
- `intl_born`: Equals 1 if the governor was born outside of the United States (as defined by the contemporary definition of the 50 states and territories); 0 otherwise. To signify internationally-born, if `birth_state` was labeled as “Other”, we recognize these governors as international.
- `intl_born_details`: To get more information on governors born outside the US, we search the bio text of governors with `birth_state` equivalent to “Other.” We extract their likely birthplace as a string.

- race/ethnicity: We use Wikipedia to source a list of all minoritized governors for the states. Then we go through our governors from the territories, and manually checked their race/ethnicity from their Bios and Wikipedia page. We broadly group into the following major categories: White, Latino/Latine/Latinx, Asian American and Pacific Islander, African American, Native American.

The data reveal notable shifts in the backgrounds of U.S. governors over the past 250 years. In the late 18th century, governors were disproportionately elite, with slightly under half having attended college. For reference, by 1870 only [one percent of the US population attended college](#), so 50 percent is exceptionally high almost a century earlier. During the Progressive Era in the early 20th century, reformers emphasized education and professionalization in public life (including in city and local government and the federal government). This is reflected in a sharp increase in college attendance among governors between 1875-1900 and 1900-1925. Note: All line plots in this essay use grey dots to represent yearly average and purple dots to represent 25 year binned averages.

```
# Note on installation: https://statsandr.com/blog/an-efficient-way-to-install-and-load-r-packages

# install.packages("rmarkdown")

# too big
# library(tidyverse)
library(dplyr)
library(tidyr)
library(stringr)
library(forcats)
library(ggplot2)
library(purrr)

# Load National Park Visitation data
gov_dataset <- read.csv("https://raw.githubusercontent.com/melaniewalsh/responsible-datasets/master/national_park_visitation.csv")

## Create long form version of the data where each year has the set of governors who governed
gov_long <- gov_dataset %>%
  # Split multiple ranges into separate rows
  separate_rows(years_in_office, sep = "\\s(?:\\d{4}\\s-\\s\\d{4}|\\d{4}\\s-\\s\\d{4})") %>%
  # Extract start and end years
  mutate(
    start = as.numeric(str_extract(years_in_office, "\\d{4}")),
    end   = as.numeric(str_extract(years_in_office, "\\d{4}$"))
  ) %>%
  # Expand each range into all years served
  rowwise() %>%
```

```

mutate(year = list(seq(start, end))) %>%
unnest(year) %>%
ungroup() %>%
filter(year >= 1775, year <= 2025) %>%
distinct()


## Let's get average college attendance rate by elected governors for each year
yearly_college <- gov_long %>%
  group_by(year) %>%
  summarise(
    total = n(),
    college = sum(college_attendance == 1, na.rm = TRUE),
    pct_college = 100 * college / total,
    .groups = "drop"
  )

# 2. Bin years into 25-year chunks
binned_college <- yearly_college %>%
  mutate(bin = cut(year, breaks = seq(1775, 2025, by = 25), include.lowest = TRUE, right =
  group_by(bin) %>%
  summarise(
    first_year = min(year), # midpoint anchor
    pct_college = mean(pct_college, na.rm = TRUE),
    .groups = "drop"
  )

# 3. Plot
ggplot() +
  geom_point(
    data = yearly_college,
    aes(x = year, y = pct_college),
    color = "gray50", alpha = 0.4
  ) +
  geom_point(
    data = binned_college,
    aes(x = first_year + 12.5, y = pct_college), # put point in middle of 25-year bin
    color = "#6A0DAD", size = 2
  ) +
  geom_line(
    data = binned_college,
    aes(x = first_year + 12.5, y = pct_college),
    color = "#6A0DAD"
  )

```

```
) +
  theme_bw() +
  xlab("") +
  ggtitle("Have elected governors become more college educated?") +
  ylab("Percent College") +
  theme(plot.title = element_text(size = 8)) +
  scale_x_continuous(breaks = seq(1775, 2025, by = 25)) # Add x-axis ticks every 25 years
```



Where Did The Data Come From? Who Collected It?

The data set came from the National Governors Association and was collected by the principal investigator Kenneth Scheve (University of Notre Dame), Sydney White (Columbia University) and Theo Serlin (King's College, Longon) as part of a larger research team at Yale University. Our sponsors included the [Yale University Department of Political Science](#), [Center for the Study of American Politics \(CSAP\)](#) and the [Yale Jackson School of Global Affairs](#).

Why Was The Data Collected? How Is The Data Used?

The data were collected to systematically document the biographical backgrounds of U.S. governors over time. This project builds on a growing literature in political science that treats the personal characteristics of politicians— such as elite education, professional background,

or family ties to politics—as potentially consequential for how they govern. Dal Bó et al. (2009) examine political dynasties in the U.S. Congress, using biographical records to identify kinship ties and measure their persistence. In the U.K., Fresh (2024) constructs a new dataset on Members of Parliament and applies natural language processing to biographical texts to code members’ business backgrounds, distinguishing established economic elites from emerging commercial elites in early modern Britain.

Compared to US Congress and the UK Parliament, U.S. governors remain relatively understudied. Yet governors have long played pivotal roles in American politics, from championing progressive reforms in the early 20th century (Berman 2019), to managing managing public health and disaster responses (see literature (Baccini and Brodeur 2021) related to COVID-19), to shaping national outcomes through gerrymandering as recently as [August 2025](#). Existing research on gubernatorial behavior, particularly during the pandemic, tends to focus narrowly on party affiliation and lacks both biographical depth and historical perspective. Sobel and Raimo (1978) is an exception and a major work on gubernatorial biographies but has not been extended to our knowledge or arranged into tabular form. This dataset helps fill that gap by systematically documenting governors’ backgrounds. This will enable new analyses of who governs the states, how elite their backgrounds are, and how such backgrounds have evolved over time.

For example, the next figure shows the share of governors with military service. As expected, service rates were higher in the decades following major conflicts such as the Civil War, World Wars I and II, Korea, and Vietnam. This pattern likely reflects the political entry of veterans after major wars. The effect is less clear after Iraq and Afghanistan, though it may simply not have fully materialized yet.

```
## Let's get average military rate by elected governors for each year
yearly_military <- gov_long %>%
  group_by(year) %>%
  summarise(
    total = n(),
    military = sum(military_service == 1, na.rm = TRUE),
    pct_military = 100 * military / total,
    .groups = "drop"
  )


# 2. Bin years into 25-year chunks
binned_military <- yearly_military %>%
  mutate(bin = cut(year, breaks = seq(1775, 2025, by = 25), include.lowest = TRUE, right =
  group_by(bin) %>%
  summarise(
    first_year = min(year),
    pct_military = mean(pct_military, na.rm = TRUE),
    .groups = "drop"
```

```

)

# 3. Plot
ggplot() +
  geom_point(
    data = yearly_military,
    aes(x = year, y = pct_military),
    color = "gray50", alpha = 0.4
  ) +
  geom_point(
    data = binned_military,
    aes(x = first_year + 12.5, y = pct_military), # center in bin
    color = "#6A0DAD", size = 2
  ) +
  geom_line(
    data = binned_military,
    aes(x = first_year + 12.5, y = pct_military),
    color = "#6A0DAD"
  ) +
  theme_bw() +
  xlab("") +
  ggtitle("Have elected governors served in the military more or less?") +
  ylab("Percent Military Service") +
  theme(plot.title = element_text(size = 8)) +
  scale_x_continuous(breaks = seq(1775, 2025, by = 25))


```



How Was The Data Collected?

To build a new individual-level data set, we scrape information for all former United States governors from the NGA website. We also scrape biographical information for each governor.

[Back to former California governors](#)



California

Gov. Arnold Schwarzenegger

Terms November 17, 2003 - January 3, 2011	Party Republican
Born July 30, 1947	Birth State Other
School University of Wisconsin, Superior	Family Married Maria Owings Shriver; four children

Figure 1: Arnold Schwarzenegger's NGA page

We use basic text extraction techniques to identify certain key biographical attributes from these biographies, such as their birth country, military service, or any legal background.

Arnold Schwarzenegger came to the governorship of California following a distinguished career in body building, business, and entertainment. **Schwarzenegger was born in Austria** and at 20 became the youngest person to win the Mr. Universe title. He won an unprecedented 12 more world bodybuilding titles. **Schwarzenegger earned a college degree from the University of Wisconsin** and became a U.S. citizen in 1983. Former President George H.W. Bush appointed him Chair of the President’s Council on Physical Fitness and Sports in 1990. He also served as Chair of the California Governor’s Council on Physical Fitness and Sports under California Governor Pete Wilson. He married Maria Shriver in 1986 and has remained closely involved in Special Olympics, an organization founded by her late mother, Eunice Kennedy Shriver.

In total, the full sample is of 2,486 individuals from 55 states and territories, including governors from Puerto Rico (1949-present), American Samoa, Guam, the Northern Mariana Islands, and the US Virgin Islands. The latter territories do not have full federal representation or rights.

The following table shows a sample portion of the data frame:

```
head(gov_dataset[c("state_territory", "governor", "party", "first_year", "school", "birth_state_territory")])
```

state_territory	governor	party	first_year	school	birth_state_territory
Alabama	Kay Ivey	Republican	2017	University of Auburn	Alabama
Alaska	Mike Dunleavy	Republican	2018	Misericordia University; University of Alaska Fairbanks	Pennsylvania
American Samoa	Pula’ali’i Nikolao Pula	Republican	2025	Menlo College; Brigham Young University; George Mason University	American Samoa
Arizona	Katie Hobbs	Democratic	2023	Northern Arizona University, Arizona State University	Arizona
Arkansas	Sarah Huckabee Sanders	Republican	2023	Ouachita Baptist University	Arkansas
California	Gavin Newsom	Democratic	2019	Santa Clara University	California

Uncertainty in the Data

The **lawyer** indicator is likely among the more error-prone variables in the dataset. It is constructed first by considering whether the governor listed his/her law school in the “school” listing, as was occasionally done. Next, we use regular expressions to identify whether a governor’s biography contains terms such as **lawyer**, **law degree**, **studied law**, **legal practice**, or **school of law**. This approach may miss individuals who practiced law but whose biographies do not contain these specific phrases, particularly in earlier periods.

Another interesting inconsistency in the data is the **birth__date** column. For example, let’s take a look at [Gov. Sam Houston Jones](#) of Louisiana.

Back to former Louisiana governors



Louisiana

Gov. Sam Houston Jones

Terms May 14, 1940 - May 9, 1944	Party Democratic
Born July 15, 1987	Passed February 8, 1978
Birth State Louisiana	School Louisiana State University
Family Married Louise Gambrell Boyer; two children	Military Service Army

Figure 2: Sam Houston Jones’ NGA page where his birth year is after his term start

Despite having his term start in 1940, the NGA site has Jones’ birth date in 1987, which—unless Jones was a time traveler— isn’t passing a basic sanity check. Obviously, the NGA made a mistake. Jones was born in 1897. But Jones is not the only one. We found 17 governors in the dataset whose birth date was **after** their term start. We manually corrected the subset whose birth date was after or too close to their term start in the dataset, but those are unlikely to be all the errors.

We extract the **school** column based on a similar method, using the **school** tag on the NGA site. However, some governors like [Kenneth Hood Mackay](#) attended college, but the NGA doesn’t list this information. We did not attempt to locate the set of governors without NGA-provided school information who may have indeed attended college, so we are likely underestimating the number of governors who attended college.

[Back to former Florida governors](#)



Florida

Gov. Kenneth Hood Mackay

Terms

December 12, 1998 - January 5, 1999

Party

Democratic

Born

March 22, 1933

Birth State

Florida

Family

Married Anne Selph; four children

National Office(s) Served

Representative

About

KENNETH H. (Buddy) MACKAY, JR., the forty-second governor of Florida, was born in Ocala, Florida on March 22, 1933. His education was attained at Ocala High School, and at the University of Florida, where he earned an undergraduate degree in 1954 and a law degree in 1961. MacKay entered politics in 1968, serving as a member of the Florida House of Representatives, a position he held until 1974. He also served as a member of the Florida State Senate from 1974 to 1980, was a member of the U.S. House of Representatives from

Current Florida Governor

 [Ron DeSantis](#)

Figure 3: Kenneth Hood Mackay's NGA page where he has no school listed in the table.

The **birth_date** and **school** columns are instructive examples of how even reliable websites backed with institutional authority can make mistakes with data entry. Those data entry mistakes can propagate into resources like this, whose quality is only ever as good as the underlying source.

Data Visualization

Here, we show some descriptive statistics. We plot the underlying yearly data, along with binned averages from each 25-year period, plotted in purple.

Thanks to this data, we can now visualize the rise in female governors in the last half-century. Historically, women have rarely been represented in the executive branch of government. The first female governor, Nellie Taylor Ross, was elected to office in Wyoming in 1925. To date, approximately 50 female governors have held office.



Figure 4: Nellie Taylor Ross was elected the first female US governor (of Wyoming) in 1925.

```
## Let's get average % female governors every year
yearly_female <- gov_long %>%
  group_by(year) %>%
  summarise(
    total = n(),
    female = sum(gender == "female", na.rm = TRUE),
    pct_female = 100 * female / total,
    .groups = "drop"
  )


# 2. Bin years into 25-year chunks
binned_female <- yearly_female %>%
  mutate(bin = cut(year, breaks = seq(1775, 2025, by = 25), include.lowest = TRUE, right = FALSE))
  group_by(bin) %>%
  summarise(
    first_year = min(year),
    pct_female = mean(pct_female, na.rm = TRUE),
  )
```

```

    .groups = "drop"
  )

# 3. Plot
ggplot() +
  geom_point(
    data = yearly_female,
    aes(x = year, y = pct_female),
    color = "gray50", alpha = 0.4
  ) +
  geom_point(
    data = binned_female,
    aes(x = first_year + 12.5, y = pct_female), # center in bin
    color = "#6A0DAD", size = 2
  ) +
  geom_line(
    data = binned_female,
    aes(x = first_year + 12.5, y = pct_female),
    color = "#6A0DAD"
  ) +
  theme_bw() +
  xlab("") +
  ggtitle("Have elected governors become more gender diverse?") +
  ylab("Percent Female Governors") +
  theme(plot.title = element_text(size = 8)) +
  scale_x_continuous(breaks = seq(1775, 2025, by = 25)) +
  scale_y_continuous(limits = c(0, 100))

```

The percentage of Ivy League-educated governors has slightly decreased over time, with approximately 25 percent of governors in each period being educated at an Ivy League institution. In previous work on political backgrounds, such as (Dal Bó et al. 2009) an Ivy League education is commonly used as a marker of elite status. This definition omits other highly selective institutions—such as Stanford, the University of Chicago, or Georgetown—which also serve as pipelines to elite political careers.

```
## Let's look at Ivy League rate
yearly_ivy <- gov_long %>%
  group_by(year) %>%
  summarise(
    total = n(),
    ivy = sum(ivy_attendance == 1, na.rm = TRUE),
    pct_ivy = 100 * ivy / total,
    .groups = "drop"
  )


# 2. Bin years into 25-year chunks
binned_ivy <- yearly_ivy %>%
  mutate(bin = cut(year, breaks = seq(1775, 2025, by = 25), include.lowest = TRUE, right =
  group_by(bin) %>%
  summarise(
    first_year = min(year),
```

```

    pct_ivy = mean(pct_ivy, na.rm = TRUE),
    .groups = "drop"
  )

# 3. Plot
ggplot() +
  geom_point(
    data = yearly_ivy,
    aes(x = year, y = pct_ivy),
    color = "gray50", alpha = 0.4
  ) +
  geom_point(
    data = binned_ivy,
    aes(x = first_year + 12.5, y = pct_ivy), # center in bin
    color = "#6A0DAD", size = 2
  ) +
  geom_line(
    data = binned_ivy,
    aes(x = first_year + 12.5, y = pct_ivy),
    color = "#6A0DAD"
  ) +
  theme_bw() +
  xlab("") +
  ggtitle("Have elected governors become more Ivy-educated?") +
  ylab("Percent with Ivy Attendance") +
  theme(plot.title = element_text(size = 8)) +
  scale_x_continuous(breaks = seq(1775, 2025, by = 25)) +
  scale_y_continuous(limits = c(0, 100))

```



Have gubernatorial offices become more professionalized over the course of American history? We track the share of governors with law degrees as a simple proxy for how “professionalized” the office is. Law has long been a standard path into politics—lawyers gain both legal expertise and entry into elite social networks. The share rises through the mid-19th century and then plateaus. Since roughly 2000, it declines modestly, consistent with more governors coming from business and other non-legal fields.

```
## Let's look at rate of governors with law degrees
yearly_lawyer <- gov_long %>%
  group_by(year) %>%
  summarise(
    total = n(),
    lawyer = sum(lawyer == 1, na.rm = TRUE),
    pct_lawyer = 100 * lawyer / total,
    .groups = "drop"
  )


# 2. Bin years into 25-year chunks
binned_lawyer <- yearly_lawyer %>%
  mutate(bin = cut(year, breaks = seq(1775, 2025, by = 25), include.lowest = TRUE, right =
  group_by(bin) %>%
  summarise(
    first_year = min(year),
```

```

    pct_lawyer = mean(pct_lawyer, na.rm = TRUE),
    .groups = "drop"
  )

# 3. Plot
ggplot() +
  geom_point(
    data = yearly_lawyer,
    aes(x = year, y = pct_lawyer),
    color = "gray50", alpha = 0.4
  ) +
  geom_point(
    data = binned_lawyer,
    aes(x = first_year + 12.5, y = pct_lawyer), # center in bin
    color = "#6A0DAD", size = 2
  ) +
  geom_line(
    data = binned_lawyer,
    aes(x = first_year + 12.5, y = pct_lawyer),
    color = "#6A0DAD"
  ) +
  theme_bw() +
  xlab("") +
  ggtitle("Have elected governors become more likely to be lawyers?") +
  ylab("Percent Lawyers") +
  theme(plot.title = element_text(size = 8)) +
  scale_x_continuous(breaks = seq(1775, 2025, by = 25)) +
  scale_y_continuous(limits = c(0, 100))

```



Which parties have produced the largest number of distinct governors? We count each person once (regardless of how many terms served), tallying their listed party affiliation. On this measure, Democrats make up the largest share, followed by Republicans, with historical parties (Whig, Democratic-Republican, Federalist, etc.) comprising smaller slices.

```
gov_dataset_deduplicated <- gov_dataset %>%
  distinct(governor, .keep_all = TRUE)


party_counts <- gov_dataset_deduplicated %>%
  transmute(party = coalesce(party, "")) %>%
  mutate(
    party = str_replace_all(party, "\\bRepublican\\b", "Republican"),
    party = str_replace_all(party, "\\bDemocrat\\b", "Democratic"),
    party = str_replace_all(party, "\\bJacksonian Democrat\\b|\\bJackson Democrat\\b", "Democratic-Republican"),
    party = str_replace_all(party, "\\bAnti-Jacksonian\\b", "National Republican"),
    party = str_replace_all(party, "\\bJeffersonian(-|)?Republican\\b", "Democratic-Republican"),
    party = str_replace_all(party, "\\bIndependent-?Republican\\b", "Republican"),
    party = str_replace_all(party, "\\s*\\(\\.\\*?\\)", ""),
    party = str_replace_all(party, "\\s+", " ")
  ) %>%
  separate_rows(party, sep = "\\s*(;|,|/|\\\\band\\\\b|&|;|\\\\s*)") %>%
  mutate(party = str_squish(party)) %>%
  filter(party != "") %>%
```

```

count(party, sort = TRUE)

party_counts_final <- party_counts %>%
  slice_max(n, n = 10) %>%
  mutate(party = fct_reorder(party, n))
ggplot(party_counts_final, aes(x = n, y = party)) +
  geom_col(fill = "#6A0DAD") +
  theme_bw() +
  theme(plot.title = element_text(size = 8)) +
  labs(
    title = "Most common party affiliations among U.S. governors",
    x = "Count", y = NULL
  ) +
  theme_minimal(base_size = 12)

```



We also want to know if governors were born in the same state that they eventually governed. We plot the share of governors who were born in their state. Surprisingly, this proportion has actually increased over time.

```

## Let's look at rate of governors with law degrees
yearly_born_in_state <- gov_long %>%
  group_by(year) %>%
  summarise(

```


```

    total = n(),
    in_state = sum(born_in_state_territory == 1, na.rm = TRUE),
    pct_in_state = 100 * in_state / total,
    .groups = "drop"
  )

# 2. Bin years into 25-year chunks
binned_born_in_state <- yearly_born_in_state %>%
  mutate(bin = cut(year, breaks = seq(1775, 2025, by = 25), include.lowest = TRUE, right =
  group_by(bin) %>%
  summarise(
    first_year = min(year),
    pct_in_state = mean(pct_in_state, na.rm = TRUE),
    .groups = "drop"
  )

# 3. Plot
ggplot() +
  geom_point(
    data = yearly_born_in_state,
    aes(x = year, y = pct_in_state),
    color = "gray50", alpha = 0.4
  ) +
  geom_point(
    data = binned_born_in_state,
    aes(x = first_year + 12.5, y = pct_in_state), # center in bin
    color = "#6A0DAD", size = 2
  ) +
  geom_line(
    data = binned_born_in_state,
    aes(x = first_year + 12.5, y = pct_in_state),
    color = "#6A0DAD"
  ) +
  theme_bw() +
  xlab("") +
  ggtitle("Percent of Governors Born in the State/Territory They Governed") +
  ylab("Percent Born In-State") +
  theme(plot.title = element_text(size = 8)) +
  scale_x_continuous(breaks = seq(1775, 2025, by = 25)) +
  scale_y_continuous(limits = c(0, 100))

```



Governors, like many politicians, have been getting older over time. In the past 25 years, the average age at entry has been about 55. This pattern likely reflects not only changes in the typical political career but also broader increases in life expectancy.

```
yearly_age <- gov_long %>%
  group_by(year) %>%
  summarise(
    avg_age = mean(age_at_start, na.rm = TRUE),
    .groups = "drop"
  )

# 2. Bin years into 25-year chunks and compute mean within each bin
binned_age <- yearly_age %>%
  mutate(bin = cut(year, breaks = seq(1775, 2025, by = 25), include.lowest = TRUE, right = FALSE))
  group_by(bin) %>%
  summarise(
    first_year = min(year),
    avg_age = mean(avg_age, na.rm = TRUE),
    .groups = "drop"
  )


# 3. Plot
ggplot() +
```



```

geom_point(
  data = yearly_age,
  aes(x = year, y = avg_age),
  color = "gray50", alpha = 0.4
) +
geom_point(
  data = binned_age,
  aes(x = first_year + 12.5, y = avg_age), # center in bin
  color = "#6A0DAD", size = 2
) +
geom_line(
  data = binned_age,
  aes(x = first_year + 12.5, y = avg_age),
  color = "#6A0DAD"
) +
theme_bw() +
xlab("") +
ggtitle("Average Age of Governors at Start of Office") +
ylab("Average Age") +
theme(plot.title = element_text(size = 8)) +
scale_x_continuous(breaks = seq(1775, 2025, by = 25)) +
scale_y_continuous(limits = c(20, 70))

```



Conclusion

The gubernatorial biography dataset gives us perspective on the range of backgrounds for important politicians in American history. Generating features for tabular data out of scraped text from the web is challenging and imperfect. But efforts like this help us notice demographic trends among US governors. We’ve pointed out a couple of interesting results around gender, professionalization, and governor age. We think future work could refine our methods and extend them to other offices, to help complete a fuller picture of the history of American government.

References

- Baccini, Leonardo, and Abel Brodeur. 2021. “Explaining Governors’ Response to the COVID-19 Pandemic in the United States.” *American Politics Research* 49 (1): 48–57. <https://doi.org/10.1177/1532673X20984458>.
- Berman, David R. 2019. *Governors and the Progressive Movement*. University Press Of Colorado.
- Dal Bó, Ernesto, Pedro Dal Bó, and Jason Snyder. 2009. “Political Dynasties.” *The Review of Economic Studies* 76 (1): 115–42.
- Fresh, Adriane. 2024. “Elite Persistence in the Era of England’s Expanding Overseas Trade.” *SSRN Electronic Journal*, ahead of print. <https://doi.org/10.2139/ssrn.4721401>.
- Sobel, Robert, and John Raimo. 1978. *Biographical Directory of the Governors of the United States, 1789-1978*. Vol. 2. Meckler Books Westport, CT.

Explore the Data

Explore the Data

Exercises

Programming Exercises

The Historical Gubernatorial Dataset is useful as a dataset for presenting interesting examples and prompting future independent exploration. Information about the governors in the dataset is well-available online, and finding historical outliers can be a fun way to engage with the history of the American Government. The exercise we present here introduces basic data

exploration skills to look at common instances in a dataframe column, and then prompts the student to explore further on their own.

Python

R